

```
1 # Dataset - https://www.kaggle.com/datasets/salader/dogs-vs-cats
```

```
1 !mkdir -p ~/.kaggle
2 !cp kaggle.json ~/.kaggle/
```

```
1 !kaggle datasets download -d saladier/dogs-vs-cats
```

```
Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600
Downloading dogs-vs-cats.zip to /content
100% 1.06G/1.06G [00:50<00:00, 23.9MB/s]
100% 1.06G/1.06G [00:50<00:00, 22.8MB/s]
```



```
1 import zipfile
2 zip_ref = zipfile.ZipFile('/content/dogs-vs-cats.zip', 'r')
3 zip_ref.extractall('/content')
4 zip_ref.close()
```

```
1 import tensorflow as tf
2 from tensorflow import keras
3 from keras import Sequential
4 from keras.layers import Dense, Conv2D, MaxPooling2D, Flatten, BatchNormalization, Dropout
```

```
1 # generators
2 train_ds = keras.utils.image_dataset_from_directory(
3     directory = '/content/train',
4     labels='inferred',
5     label_mode = 'int',
6     batch_size=32,
7     image_size=(256,256)
8 )
9
10 validation_ds = keras.utils.image_dataset_from_directory(
11     directory = '/content/test',
12     labels='inferred',
13     label_mode = 'int',
14     batch_size=32,
15     image_size=(256,256)
16 )
```

```
Found 5000 files belonging to 2 classes.
```

```
1 # Normalize
2 def process(image,label):
3     image = tf.cast(image/255. ,tf.float32)
4     return image,label
5
6 train_ds = train_ds.map(process)
7 validation_ds = validation_ds.map(process)
```

```
1 # create CNN model
2
3 model = Sequential()
4
5 model.add(Conv2D(32,kernel_size=(3,3),padding='valid',activation='relu',input_shape=(256,256,3)
6 model.add(BatchNormalization())
7 model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))
8
9 model.add(Conv2D(64,kernel_size=(3,3),padding='valid',activation='relu'))
10 model.add(BatchNormalization())
11 model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))
12
13 model.add(Conv2D(128,kernel_size=(3,3),padding='valid',activation='relu'))
14 model.add(BatchNormalization())
15 model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))
16
17 model.add(Flatten())
18
19 model.add(Dense(128,activation='relu'))
20 model.add(Dropout(0.1))
21 model.add(Dense(64,activation='relu'))
22 model.add(Dropout(0.1))
23 model.add(Dense(1,activation='sigmoid'))
```

```
1 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 254, 254, 32)	896
batch_normalization (Batch Normalization)	(None, 254, 254, 32)	128
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 125, 125, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 60, 60, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	0
flatten (Flatten)	(None, 115200)	0
dense (Dense)	(None, 128)	14745728
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 64)	8256
dropout_1 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 1)	65


=====  
Total params: 14848193 (56.64 MB)  
Trainable params: 14847745 (56.64 MB)

```
1 model.compile(optimizer='adam',loss='binary_crossentropy',metrics=['accuracy'])
```

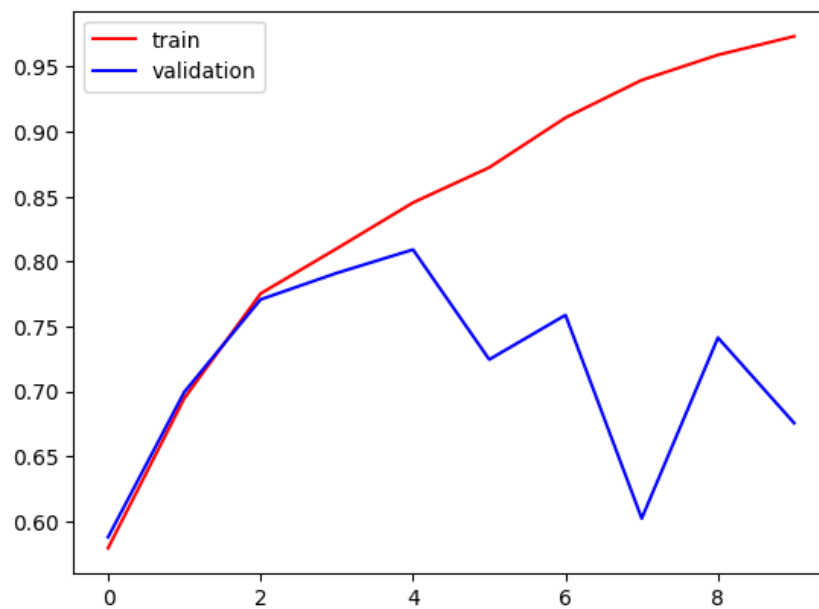
```
1 history = model.fit(train_ds,epochs=10,validation_data=validation_ds)
```

```
Epoch 1/10
625/625 [=====] - 135s 192ms/step - loss: 1.2646 - accuracy: 0.5794 - val_loss: 0.6
Epoch 2/10
625/625 [=====] - 123s 196ms/step - loss: 0.5729 - accuracy: 0.6948 - val_loss: 0.5
Epoch 3/10
625/625 [=====] - 115s 184ms/step - loss: 0.4759 - accuracy: 0.7753 - val_loss: 0.4
Epoch 4/10
625/625 [=====] - 117s 186ms/step - loss: 0.4164 - accuracy: 0.8101 - val_loss: 0.4
Epoch 5/10
625/625 [=====] - 112s 178ms/step - loss: 0.3524 - accuracy: 0.8454 - val_loss: 0.4
Epoch 6/10
625/625 [=====] - 127s 203ms/step - loss: 0.2965 - accuracy: 0.8724 - val_loss: 0.5
Epoch 7/10
625/625 [=====] - 128s 204ms/step - loss: 0.2139 - accuracy: 0.9108 - val_loss: 0.5
Epoch 8/10
625/625 [=====] - 134s 215ms/step - loss: 0.1553 - accuracy: 0.9398 - val_loss: 2.7
Epoch 9/10
625/625 [=====] - 119s 190ms/step - loss: 0.1137 - accuracy: 0.9592 - val_loss: 1.4
Epoch 10/10
625/625 [=====] - 131s 209ms/step - loss: 0.0789 - accuracy: 0.9735 - val_loss: 2.7
```

```
1 import matplotlib.pyplot as plt
2
3 plt.plot(history.history['accuracy'],color='red',label='train')
4 plt.plot(history.history['val_accuracy'],color='blue',label='validation')
5 plt.legend()
6 plt.show()
```




```
1 plt.plot(history.history['accuracy'],color='red',label='train')
2 plt.plot(history.history['val_accuracy'],color='blue',label='validation')
3 plt.legend()
4 plt.show()
```



```

1 plt.plot(history.history['loss'],color='red',label='train')
2 plt.plot(history.history['val_loss'],color='blue',label='validation')
3 plt.legend()
4 plt.show()


```



```

1 plt.plot(history.history['loss'],color='red',label='train')
2 plt.plot(history.history['val_loss'],color='blue',label='validation')
3 plt.legend()
4 plt.show()

```



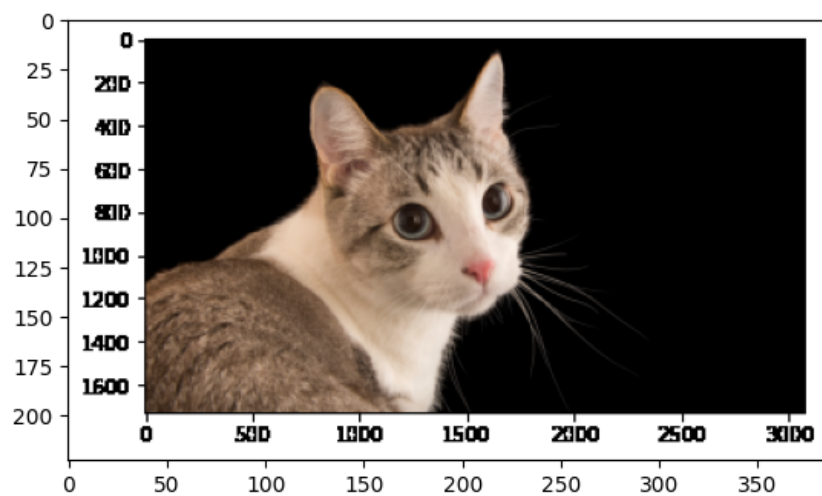
```
1 # ways to reduce overfitting
2
3 # Add more data
4 # Data Augmentation -> next video
5 # L1/L2 Regularizer
6 # Dropout
7 # Batch Norm
8 # Reduce complexity
```

```
1 import cv2
```

```
1 test_img = cv2.imread('/content/cat.png')
```

```
1 plt.imshow(test_img)
```

```
<matplotlib.image.AxesImage at 0x7de22a38a4a0>
```



```
1 test_img.shape
```

```
(223, 386, 3)
```

```
1 test_img = cv2.resize(test_img,(256,256))
```

```
1 test_input = test_img.reshape((1,256,256,3))
```

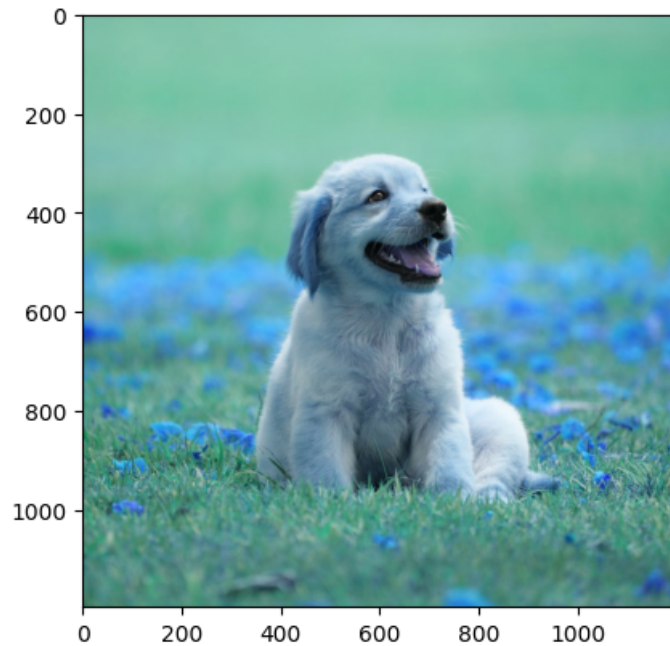
```
1 model.predict(test_input)
```

```
1/1 [=====] - 0s 388ms/step  
array([[0.]], dtype=float32)
```

```
1 test_img2 = cv2.imread('/content/dog.jpg')
```

```
2 plt.imshow(test_img2)
```

```
<matplotlib.image.AxesImage at 0x7de1f7fd8d90>
```



```
1 test_img2 = cv2.resize(test_img2,(256,256))
```

```
1 test_input2 = test_img2.reshape((1,256,256,3))
```

```
1 model.predict(test_input2)
```

```
1/1 [=====] - 0s 26ms/step  
array([[1.]], dtype=float32)
```